

# Machine Learning Final Assignment

Name: Punam Das

ID : 21-44946-2

## Q1) Locally Weighted Regression vs Linear Regression:

- **Linear Regression:** Linear regression aims to fit a linear relationship between the independent variable(s)  $X$  and the dependent variable  $y$ . The formula for linear regression is:

$$y = \beta_0 + \beta_1 X + \epsilon$$

Where:

- $y$  is the dependent variable.
- $X$  is the independent variable.
- $\beta_0$  is the intercept.
- $\beta_1$  is the slope coefficient.
- $\epsilon$  is the error term.
- **Locally Weighted Regression (LWR):** LWR differs from linear regression by giving more weight to data points closer to the point at which a prediction is being made. The formula for LWR is:

$$y = \sum_{i=1}^n w_i (y_i - \beta_0 - \beta_1 X_i)$$

Where:

- $w_i$  is the weight assigned to each data point.
- $(y_i - \beta_0 - \beta_1 X_i)$  is the weighted residual.

**Advantage of LWR over Linear Regression:** The advantage of LWR is that it can capture more complex patterns in the data by considering local information. Linear regression assumes a global linear relationship between variables, which may not always be the case in real-world scenarios.

## Q2) Binary Logistic Regression for Tumor Prediction:

Binary logistic regression is used to model the probability that a given observation belongs to a particular category (in this case, malignant or benign tumor). The logistic function is used to transform the output of linear regression into a probability between 0 and 1. The model output  $y$  is:

$$y = \frac{1}{1 + e^{-z}}$$

Where  $z = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n$  is the linear combination of features and coefficients.

The learning algorithm for binary logistic regression involves minimizing the logistic loss function using techniques like gradient descent. The logistic loss function is:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

Where:

- $h_{\theta}(x)$  is the logistic function.
- $y^{(i)}$  is the actual label of the  $i$ -th observation.
- $x^{(i)}$  is the feature vector of the  $i$ -th observation.

Once trained, the logistic regression model can predict the tumor type (malignant or benign) for a new patient by calculating the probability using the logistic function and applying a threshold (e.g., 0.5) to classify the tumor.

## Q3.a) Softmax Regression:

- **Cost Function  $J(w)$ :** For  $nn$  training items, the softmax cost function is:

$$J(w) = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K y_k^{(i)} \log(\hat{y}_k^{(i)})$$

Where:

- $y_k(i)$  is the  $k$ -th component of the one-hot encoded label for the  $i$ -th item.
- $\hat{y}_k(i)$  is the predicted probability of class  $k$  for the  $i$ -th item.

- **Softmax Output Function  $f(x;w)$ :** The softmax output function for  $K$  classes is:

$$f(x;w) = \frac{e^{w_j^T x}}{\sum_{k=1}^K e^{w_k^T x}}$$

### Q3.b) Relationship between Softmax and Binary Logistic Regression:

Softmax regression is a generalization of binary logistic regression to multiple classes. In binary logistic regression, there are only two possible outcomes (0 or 1), while softmax regression deals with multiple mutually exclusive classes.

### Q4) Ridge/L2 Regularization Penalty Term:

The penalty term of ridge/L2 regularization is the sum of squares of all the model coefficients (except the intercept term) multiplied by the regularization parameter  $\lambda$ . It is added to the cost function to prevent overfitting. The formula for the ridge regularization penalty term is:

$$\text{Penalty Term} = \lambda \sum_{j=1}^p \beta_j^2$$

Where:

- $\lambda$  is the regularization parameter.
- $\beta_j$  are the coefficients of the model.

This penalty term reduces overfitting by shrinking the coefficients towards zero, thereby reducing the model's complexity.

### Q5.a) Policy Iteration for MDP:

Initialize the value function  $V(s)$  arbitrarily

Repeat until convergence {

Policy Evaluation:

for each state  $s$  {

$$V(s) = \sum_a \pi(a|s) \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma V(s')]$$

}

Policy Improvement:

for each state  $s$  {

$$\pi(s) = \operatorname{argmax}_a \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma V(s')]$$

}

}

### Q5.b) Advantage of Exploration-based Policy ( $\epsilon$ -Greedy):

The advantage of using an exploration-based policy like  $\epsilon$ -greedy to solve an MDP is that it allows the agent to explore different actions and learn about the environment. By occasionally choosing random actions (with probability  $\epsilon$ ), the agent can discover potentially better strategies and avoid getting stuck in suboptimal policies.

#### **Q6.a) Q-Learning as an Off-Policy Algorithm:**

Q-learning is an off-policy algorithm because it learns the value of the optimal policy while following a different (possibly exploratory) policy. It updates the Q-values based on the maximum expected future reward, regardless of the action actually taken.

#### **Q6.b) Difference between On-Policy and Off-Policy Algorithms:**

- **On-Policy Algorithms:** On-policy algorithms update the policy being followed directly based on the observed experience. They evaluate or improve the same policy that is used to interact with the environment (e.g., SARSA).
- **Off-Policy Algorithms:** Off-policy algorithms learn the value function or optimal policy from data collected using a different policy. They can learn from historical data or explore alternative strategies while still aiming to find the optimal policy (e.g., Q-learning).